

Azure SQL - Basics

Agenda

- What is Azure SQL?
- Deployment Models
- Setup - SQL Server, Azure SQL Database, Firewall
- Connect through Azure Data Studio
- Database Security Features

Azure SQL

- Fully-managed, database services in Azure
- Create only databases in the cloud
- Provides latest stable release of SQL Server Database
- Scalable and highly available
- Supports geographic replication and failover

Deployment Models

- **SQL Server on Azure VM**
 - Full control on SQL Server, databases, and underlying infrastructure
 - Deployed inside Virtual Network
- **Azure SQL databases**
 - Most features of on-prem SQL Server are supported (some are not)
 - Cheaper, and highly-scalable
 - Supports SQL Server, MySQL and PostgreSQL databases
- **Azure SQL Managed Instance**
 - Nearly compatible with on-prem SQL Server
 - Managed service; deployed inside Virtual Network

Azure SQL - Advanced

Agenda

- Deployment Models
- Purchasing Models
- Resource Sharing Options
- Service Tiers
- Security Features
- Geo replication & Failover groups
- High Availability

Azure SQL

- Fully-managed, database services in Azure
- Create only databases in the cloud
- Provides latest stable release of SQL Server Database
- Scalable and highly available
- Supports geographic replication and failover

Deployment Models

- **SQL Server on Azure VM**
 - Full control on SQL Server, databases, and underlying infrastructure
 - Deployed inside Virtual Network
- **Azure SQL databases**
 - Most features of on-prem SQL Server are supported (some are not)
 - Cheaper, and highly-scalable
 - Supports SQL Server, MySQL and PostgreSQL databases
- **Azure SQL Managed Instance**
 - Nearly compatible with on-prem SQL Server
 - Managed service; deployed inside Virtual Network

Purchasing Models

- **DTU based Model**

- DTU stands for Database Transaction Unit
- Combination of compute, storage and IOPS
- Pre-configured resource option - Pay only for DTUs

- **vCore based Model**

- Select & scale independently - compute, storage and IOPS
- More flexible and transparent
- Azure Hybrid Benefit – use existing licenses to save cost

Resource Sharing Options

- Only available in Azure SQL database
- Available for both DTU and vCore purchase models
- **Single Database**
 - Each database gets its own set of dedicated resources
- **Elastic Pool of Databases**
 - Creates a pool of resources (compute, IOPS etc.)
 - Multiple databases can run on the same pool, sharing the resources

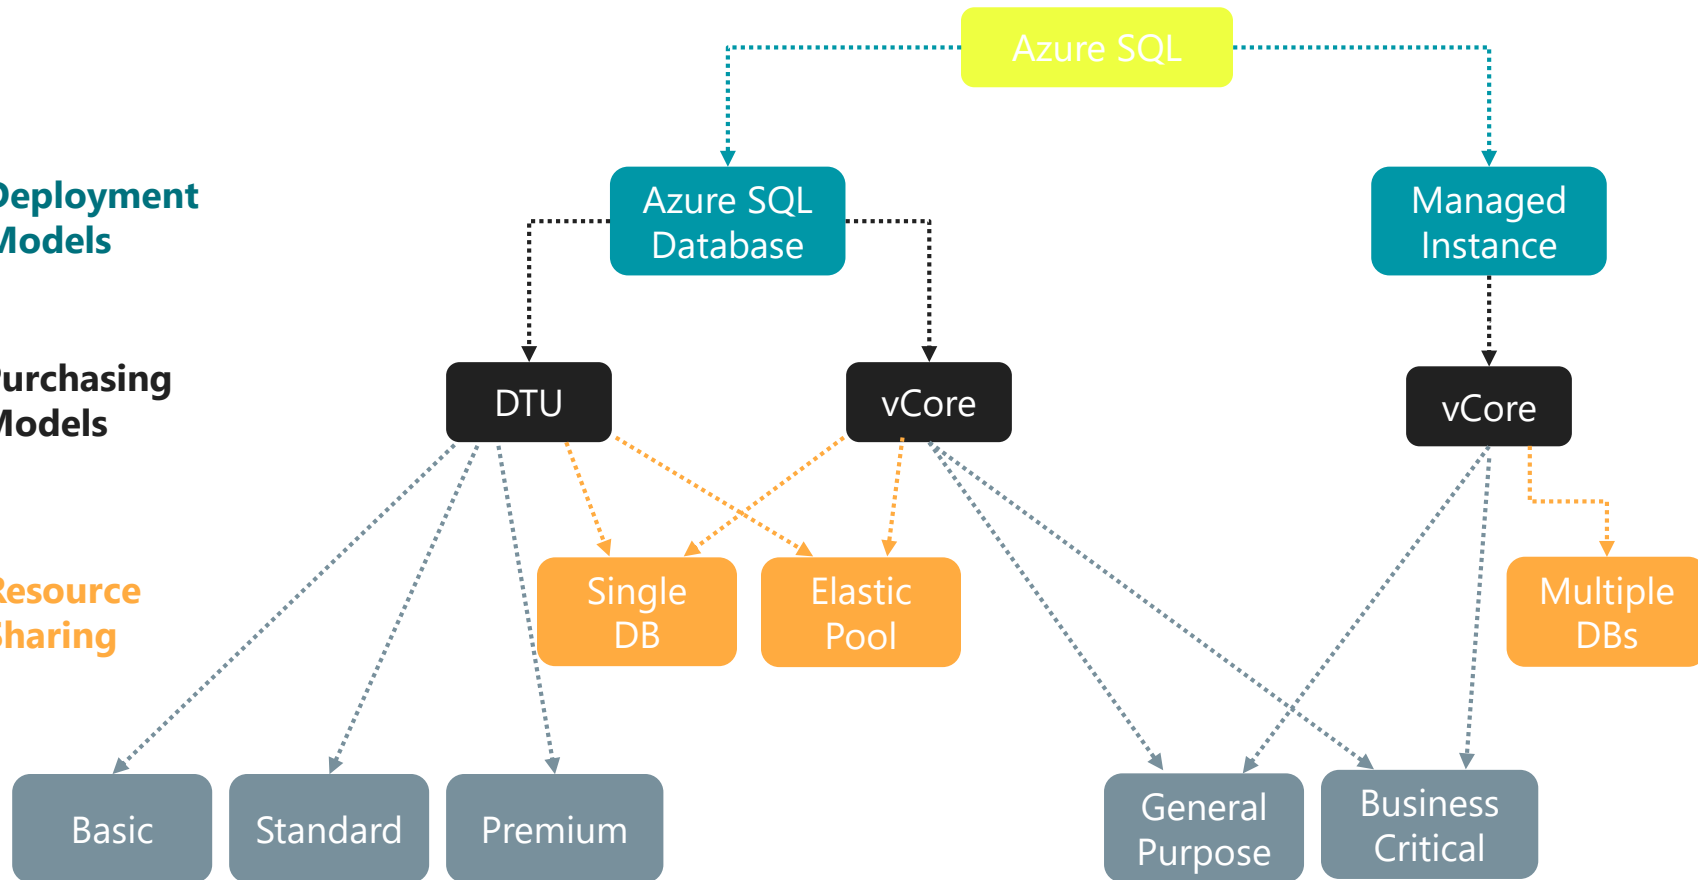
Service Tiers

- **DTU model**
 - Basic
 - Standard
 - Premium
- **vCore model**
 - General purpose
 - Business critical
 - Hyperscale

**Deployment
Models**

**Purchasing
Models**

**Resource
Sharing**



High Availability

- Standard Availability Model
- Premium Availability Model

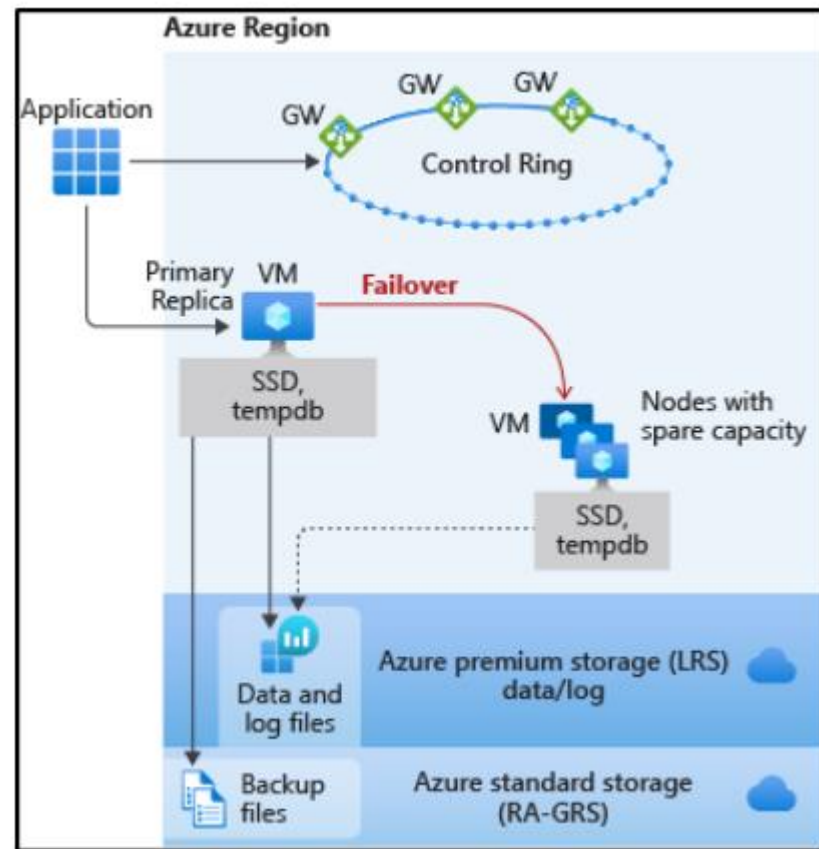
Standard Availability Model

Stateless Compute Layer

- Runs in a VM
- Only contains cached & transient data in SSD
- On crash, auto failover is performed by using spare capacity

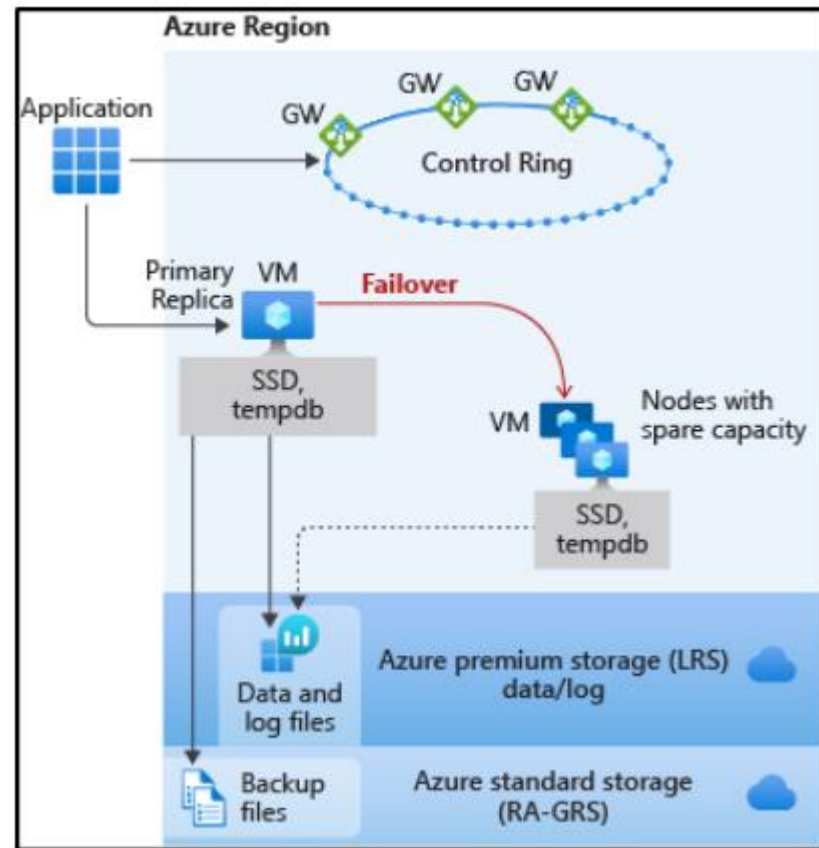
Stateful Data Layer

- Data is stored in Premium Storage (mdf/ldf files)
- Data is safe even if machine with SQL engine crashes



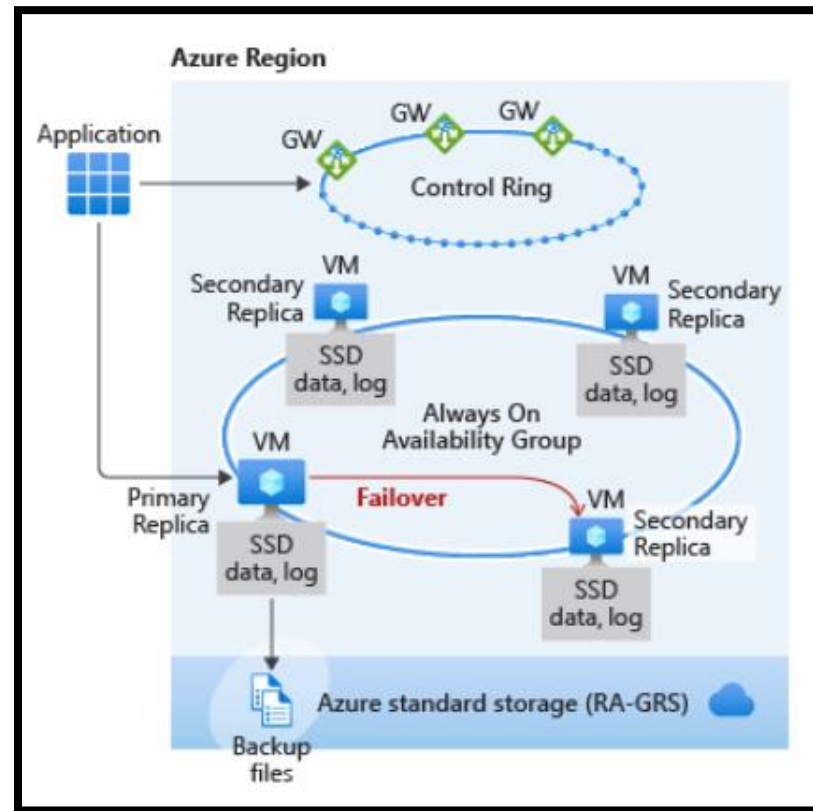
Standard Availability Model

- Cheaper option
- Less performant than Premium availability model
- Service Tiers supported
 - Basic
 - Standard
 - General Purpose



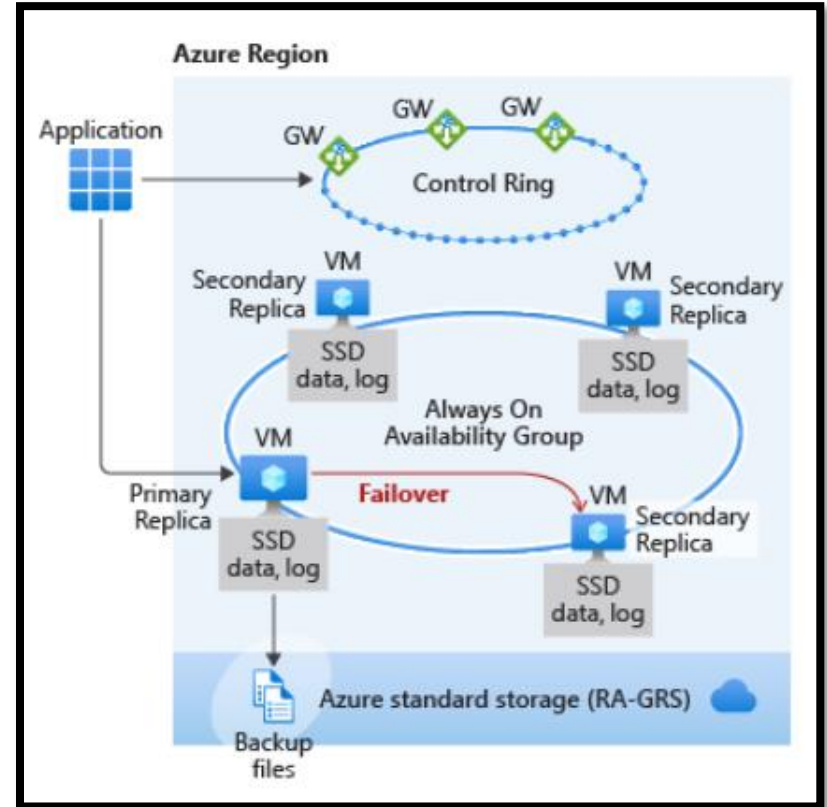
Premium Availability Model

- Compute & data are on same node
- Data is stored on attached SSDs (mdf/ldf files)
- Compute & data are replicated on multiple secondary nodes in cluster
- Uses Always On Availability Groups to provide high availability
- If primary node crashes, failover is performed

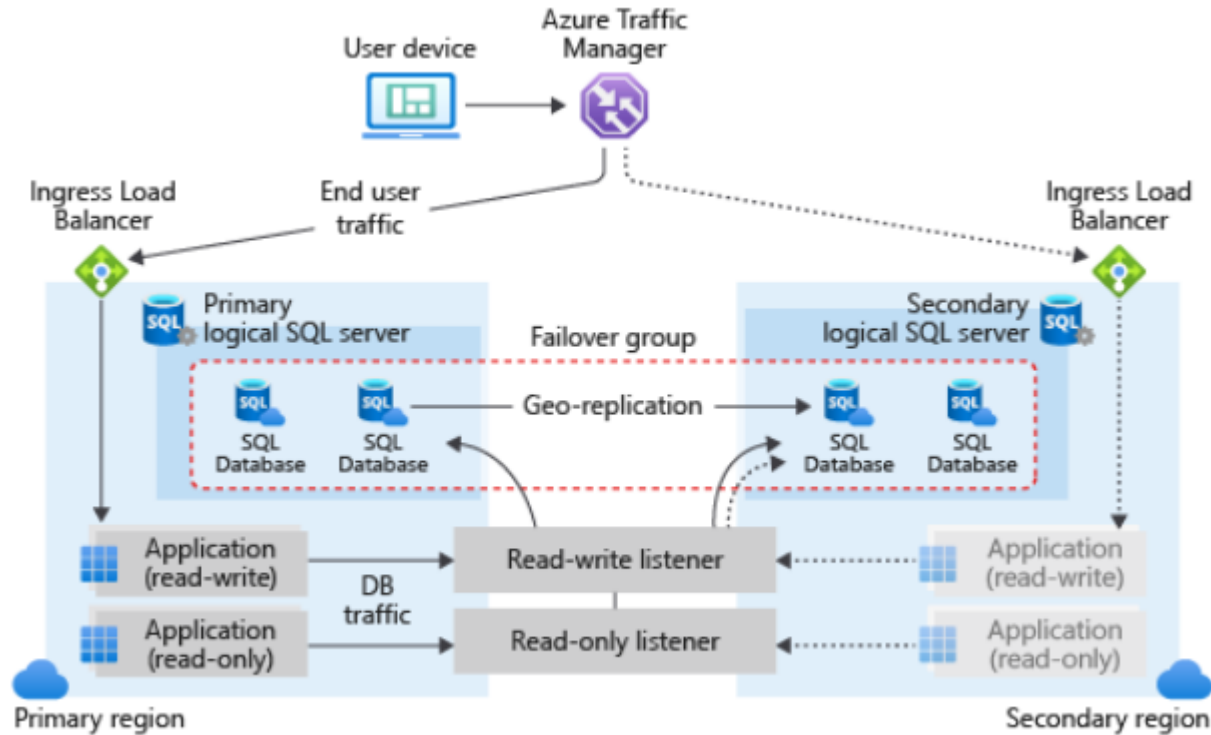


Premium Availability Model

- Expensive option
- More performant than Standard availability model
- Service Tiers supported
 - Premium
 - Business critical



Geo-Replication and Failover Groups



Azure Cosmos DB - Basics

Agenda

- NoSQL Concepts
- NoSQL Database Types
- Introduction to Azure Cosmos DB
- Cosmos DB APIs
- Resource Model
- Working with Core SQL API
- Connecting to Azure function

NoSQL Concepts

- Non-relational / non-SQL
- Schema-free / Flexible schema
- Distributed – Data is locally distributed amongst multiple nodes
- Replicated – Multiple copies of the data are created
- Horizontal scaling – Sharding across servers
- Provides
 - High Availability
 - High Throughput
 - Low Latency
- Maintaining consistency is challenging – provides eventual consistency

Use Cases

- Content management
- Personalization
- Build networks of entities
- Social media data

NoSQL Database Types

- Key-Value stores
- Document stores
- Wide-column / Column-family / Columnar stores
- Graph stores

Key-Value Store

- Data is associated with unique key
- Value is considered binary
- Query only on keys
- Optimized for simple lookups

| Key | Value |
|--------|------------------------------|
| AAAAA | 1101001111010100110101111... |
| AABAB | 1001100001011001101011110... |
| DFA766 | 0000000000101010110101010... |
| FABCC4 | 1110110110101010100101101... |

Opaque to
data store

Document Store

- Data stored is known as Document
- Document is associated with unique key
- Document has flexible schema
- Document can be different formats - XML, JSON, YAML, BSON etc.
- Query on keys or on fields inside document

| Key | Document |
|------|---|
| 1001 | <pre>{ "CustomerID": 99, "OrderItems": [{ "ProductID": 2010, "Quantity": 2, "Cost": 520 }, { "ProductID": 4365, "Quantity": 1, "Cost": 18 }], "OrderDate": "04/01/2017" }</pre> |
| 1002 | <pre>{ "CustomerID": 220, "OrderItems": [{ "ProductID": 1285, "Quantity": 1, "Cost": 120 }], "OrderDate": "05/08/2017" }</pre> |

Columnar Store

- Uses column-oriented format and denormalized approach
- Looks very similar to relational structure
- Columns are grouped into column families, which are retrieved together

- Key is stored together

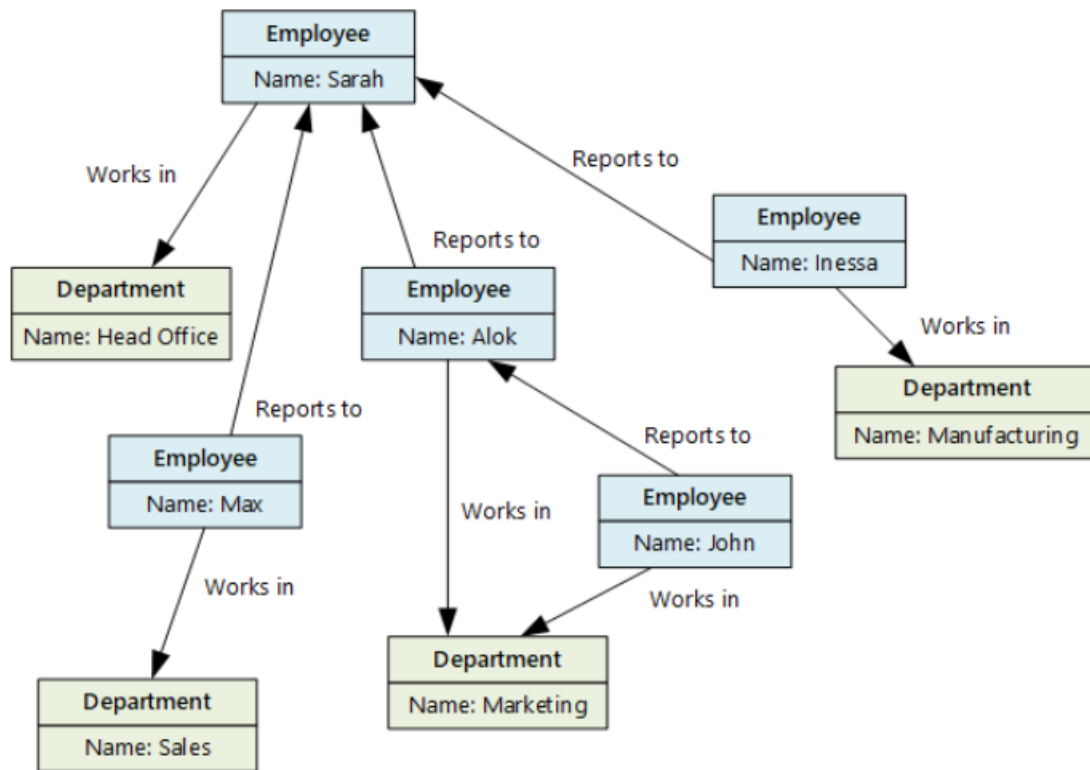
| CustomerID | Column Family: Identity |
|------------|--|
| 001 | First name: Mu Bae Last name: Min |
| 002 | First name: Francisco Last name: Vila Nova Suffix: Jr. |
| 003 | First name: Lena Last name: Adamczyk Title: Dr. |

| CustomerID | Column Family: Contact Info |
|------------|--|
| 001 | Phone number: 555-0100 Email: someone@example.com |
| 002 | Email: vilanova@contoso.com |
| 003 | Phone number: 555-0120 |

' stored

Graph Store

- Based on nodes and edges
- Maintains relationship between entities
- Queries traverses nodes by using relationships



Azure Cosmos DB

- Fully managed NoSQL platform
- Supports typical NoSQL features
 - Distribution, replication, horizontal scaling, high throughput, low latency etc.
- Multi-API support
 - Create any type of NoSQL store
 - Supports five APIs
- Supports global distribution
 - Multiple read regions support with one write region
 - Multiple write regions support
- Automatic expiration of documents using Time-to-Live feature

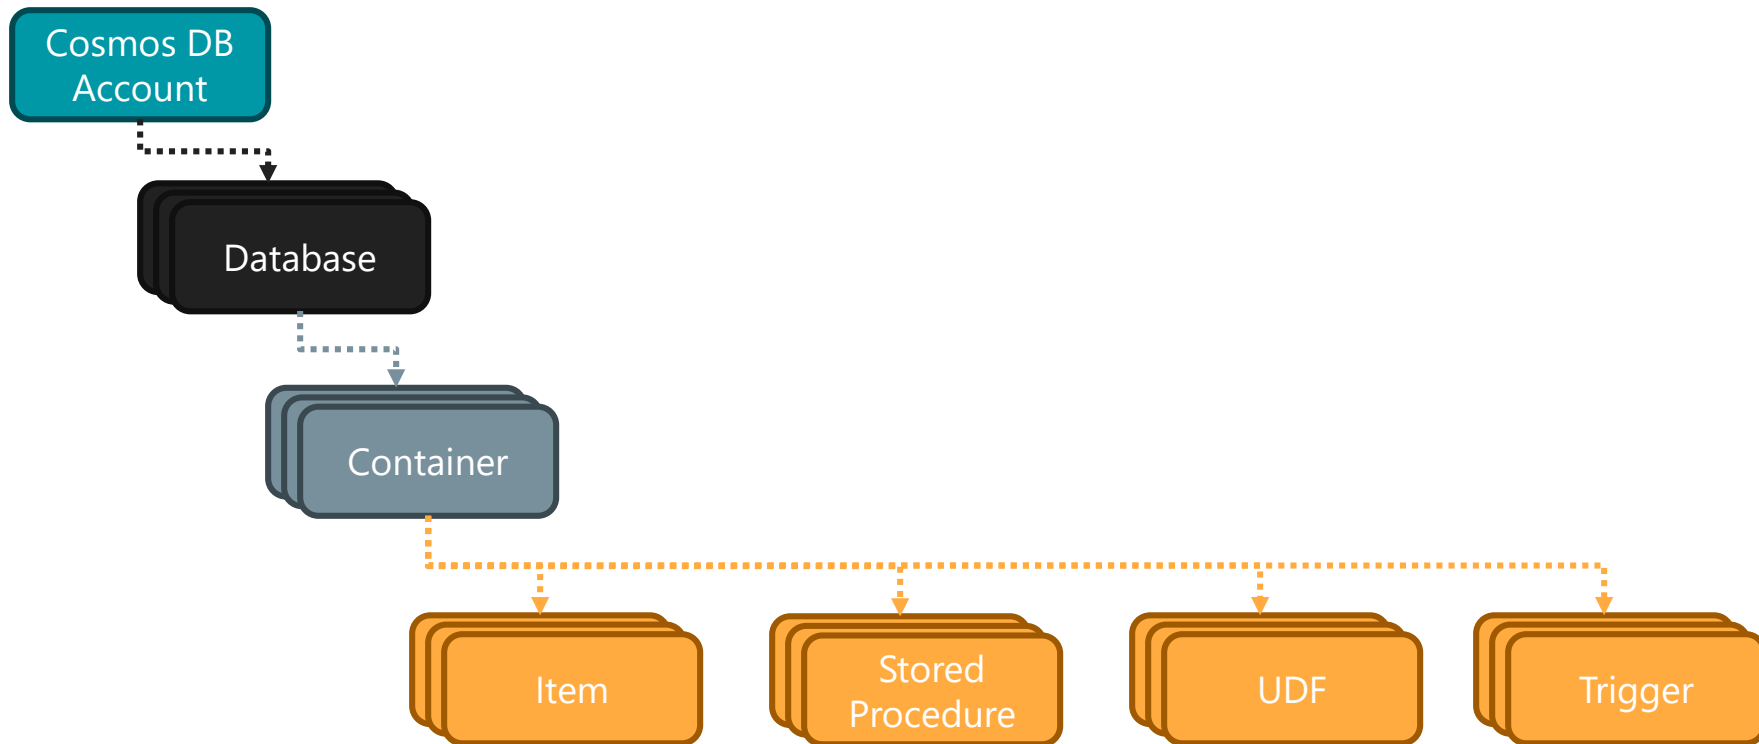
Cosmos DB APIs

- **Table API**
 - Key-value store
 - Similar to Azure Table Storage
- **Core (SQL) API**
 - Document store
 - Earlier known as Azure Document DB
 - Developed by Microsoft
 - Store data in JSON format
 - Supports SQL-like syntax to query the data
 - Build stored procedures, user-defined functions and triggers in JavaScript

Cosmos DB APIs

- **MongoDB API**
 - Document store
 - MongoDB service in Azure, powered by Cosmos DB
 - Use existing MongoDB libraries, tools and applications
- **Cassandra API**
 - Columnar store
 - Cassandra service in Azure, powered by Cosmos DB
 - Use existing Cassandra libraries, tools and applications
- **Gremlin API**
 - Graph store
 - Based on open-source Apache Gremlin

Resource Model



Resource Model

| Azure Cosmos entity | SQL API | Cassandra API | MongoDB API | Gremlin API | Table API |
|-------------------------|-----------|---------------|-------------|--------------|-----------|
| Cosmos database | Database | Keyspace | Database | Database | NA |
| Cosmos container | Container | Table | Collection | Graph | Table |
| Cosmos item | Document | Row | Document | Node or edge | Item |

Azure Cosmos DB - Advanced

Agenda

- Partitioning
- Throughput
- Global Distribution

Azure Cosmos DB

- Fully managed NoSQL platform
- Supports typical NoSQL features
 - Distribution, replication, horizontal scaling, high throughput, low latency etc.
- Multi-API support
 - Create any type of NoSQL store
 - Supports five APIs
- Supports global distribution
 - Multiple read regions support with one write region

Partitioning

- Every container has a user-defined Partition Key
- Data is stored in separate partitions, based on Partition Key
- All data with same partition key is stored in same partition
- Records can be retrieved efficiently using Partition Key
- One partition has max size of 20 GB

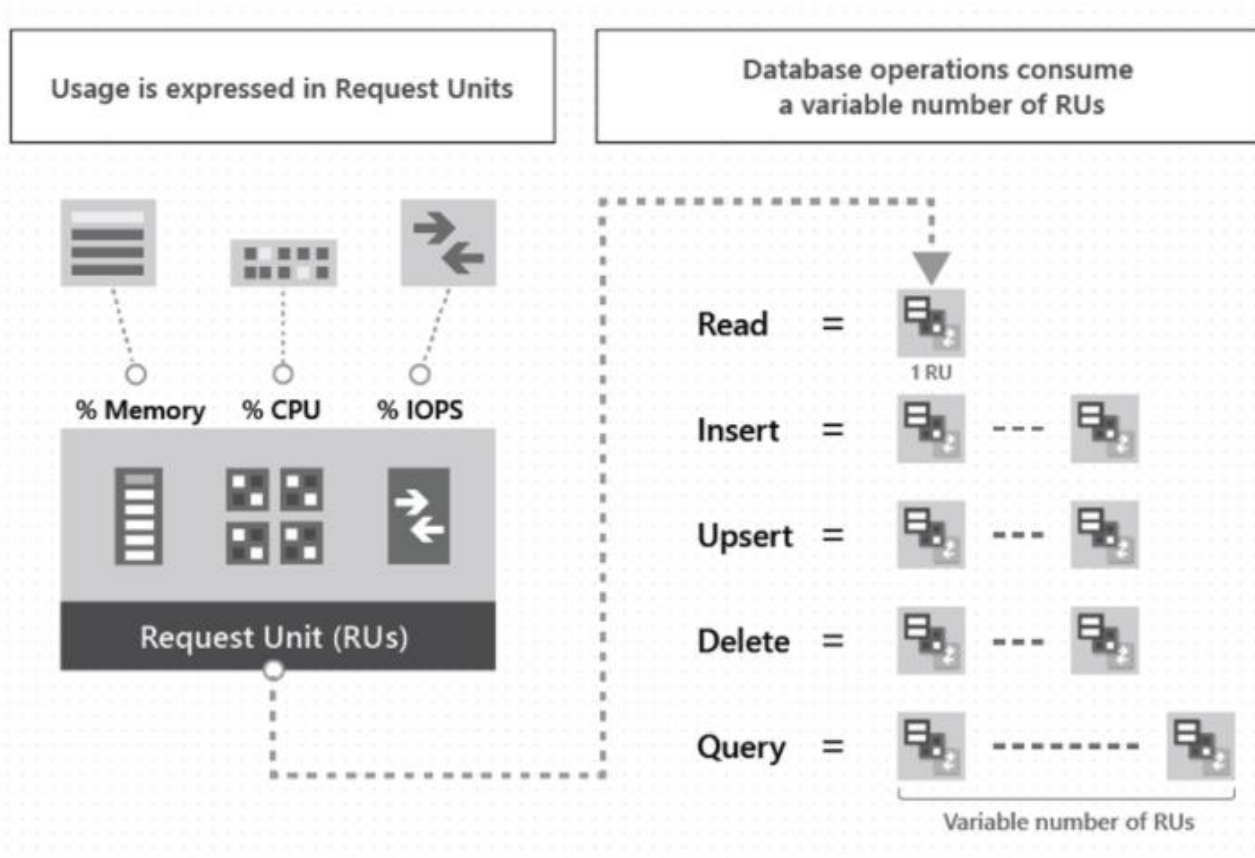
Partition Key – Design Considerations

- Partition key should have high cardinality – many different values
- Should be decided based on workload
 - Read-heavy vs write-heavy
- **Write-heavy**
 - Try to write data to as many partitions as possible
 - Choose a partition key that maximizes the parallelizability of writes
 - Avoid hot-partitions
- **Read-heavy**
 - Try to read from one partition as much as possible
 - Choose a partition key that optimizes the most common queries
 - Avoid fan-out

Throughput

- Number of requests that can be served in an instance of time
- Options
 - Provisioned throughput
 - Autoscaling throughput
 - Serverless throughput
- Throughput in Cosmos DB is defined using Request Units (RUs)
 - RUs is not equal to number of Requests
 - One read or write request can consume more than 1 RU
 - Each RU is a combination of CPU, memory and IOPS
 - Depending on type of query, RUs are consumed
- To handle more requests, provision more RUs

Throughput



Factors affecting Request Units

- Item size
- Item indexing – by default all attributes in a document are indexed
- Attributes within a document – if indexing is enabled
- How data is distributed in partitions
- Type of query
 - Read or write query
 - Searching within a document consumes higher RUs
 - Querying on partition key consumes less RUs
 - Applying filters on attributes of document increases RUs etc.

Azure Cosmos DB

SQL

SQL

{LEAF}

API for MongoDB



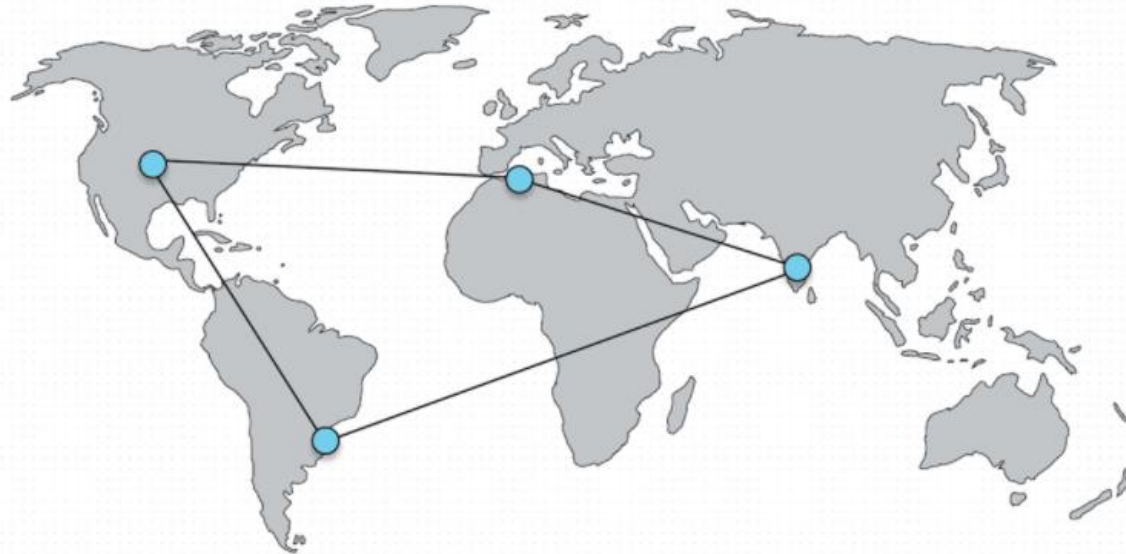
Gremlin



Cassandra



Table




Key-Value



Column-Family



Documents



Graph



Guaranteed speed at any scale

Simplified application development

Mission-critical ready

Fully managed and cost effective

Global Distribution

- Supports single read-write and multiple read regions
- Supports multiple read-write regions
- Provides global low latency, high availability and high throughput
- Handles conflict management
 - LWW (Last-Write-Wins) algorithm – depends on timestamp
 - Custom-defined algorithm